



# Day 24





# Indexes

- Structure that improves data retrieval
- Separate table to store a sorted values of one or more columns
- Index is created automatically for primary key
  - Other columns are optional
- Indexes are useful for
  - Reducing query time
  - Grouping
  - Sorting
  - Minimum and maximum values



# Index

Primary key (`prog_id`) index

Values are sorted

<code>idx_prog_id(prog_id)</code>
2
3
7
10

<code>prog_id</code>	<code>name</code>	<code>release_date</code>
10	Game of Thrones	2011-04-17
2	Get Smart	1965-09-18
7	Twin Peaks	1990-04-08
3	Hinterland	2013-10-29

`select * from tv_shows where prog_id = 3;`  An indexed column

Search the index of `prog_id`. Since the values are sorted, only 2 entries are search.  
If we have to search the `prog_id` from `tv_show` table, we will have to search 4 entries to get the result



# Creating Index

- Creating an index

```
create index idx_name  
on tv_shows (name)
```

Index's name

For this column

tv_shows	
🔑	prog_id INT(11)
💎	name CHAR(64)
💎	lang CHAR(16)
💎	official_site VARCHAR(256)
💎	rating ENUM(...)
💎	user_rating FLOAT
💎	release_date DATE
💎	image BLOB
Indexes ▶	

- Speed up all queries where the predicate involves the `name` column only

```
select * from tv_shows where name like 'Game%';
```

- Index will slow down inserts
  - Have to insert into the index at the correct (sorted) location



# Multi Column Index

```
create index lang_release_date_idx
on tv_shows (release_date, lang);
```

tv_shows	
🔑	prog_id INT(11)
💎	name CHAR(64)
💎	lang CHAR(16)
💎	official_site VARCHAR(256)
💎	rating ENUM(...)
💎	user_rating FLOAT
💎	release_date DATE
💎	image BLOB
Indexes ▶	

release_date	lang	prog_id	name	lang	release_date
2008-11-30	EN	20	Borgen	SE	2010-09-26
2015-10-10	EN	9	Wallander	EN	2008-11-30
2015-09-20	IS	25	The Last Kingdom	EN	2015-10-10
2010-09-26	SE	5	Ófærð	IS	2015-09-20

```
select * from tv_shows
where release_date > '2014-01-10' and lang like 'is'
```



# Multi Column Index

- Column with higher cardinality should be placed first
  - Index cardinality means the number of unique values
  - A unique

release\_date has a higher cardinality than lang  
Should be placed at the left most

release_date	lang
2008-11-30	EN
2015-10-10	EN
2015-09-20	IS
2010-09-26	SE

prog_id	name	lang	release_date
20	Borgen	SE	2010-09-26
9	Wallander	EN	2008-11-30
25	The Last Kingdom	EN	2015-10-10
5	Ófærð	IS	2015-09-20





# Using Multi Column Index

- Order is important

```
select * from tv_shows
  where release_date > '2014-01-01' and lang like 'is'
```

- Use of index from left to right

```
select * from tv_shows
  where release_date > '2010-01-01'
```



```
select * from tv_shows
  where lang = 'en'
```







# Using Multi Column Index

user\_rating

lang

rating

where **user\_rating** > 5

where **user\_rating** > 5 and **lang** = 'EN'

where **user\_rating** > 5 and **lang** = 'EN' and **rating** = 'NC16'

where **lang** = 'EN'

where **lang** = 'EN' and **rating** = 'NC16'

where **lang** = 'EN' and **user\_rating** > 5



# Covering Index

- Results can be returned by just querying the index
  - Index contains all the columns to fulfil the query
  - The index holds all the data

```
create index lang_release_date_idx  
on tv_shows (release_date, lang);
```



```
select distinct(lang) from tv_shows  
where release_date > '2019-01-01'
```

Result can be  
returned from query

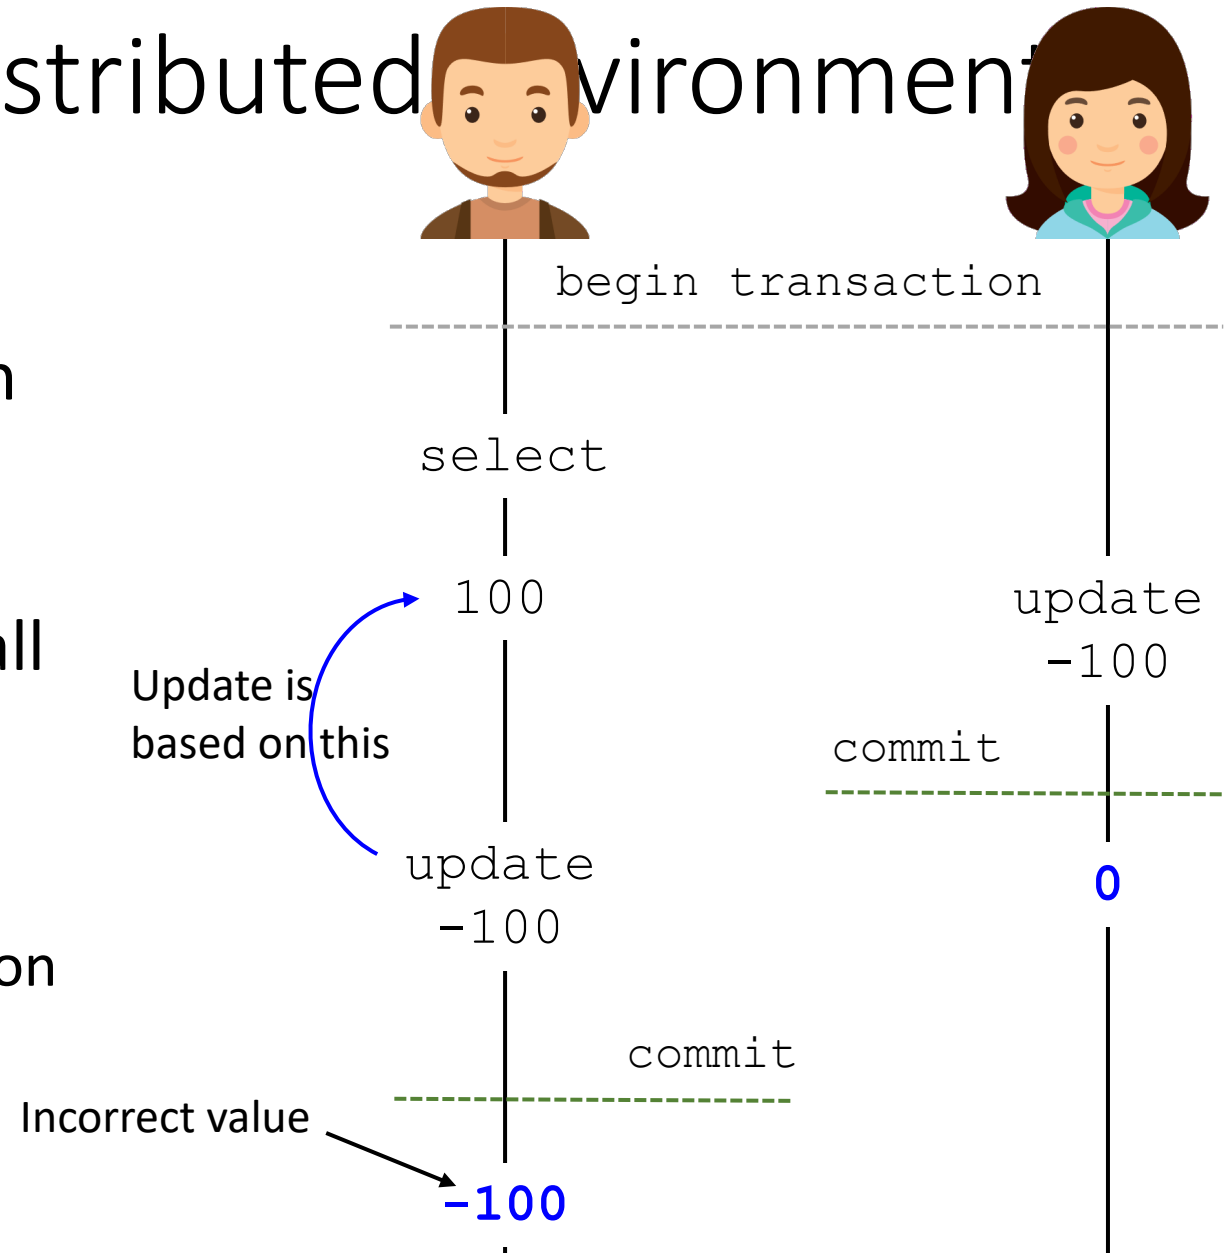


```
select release_date from tv_shows  
where lang like 'en'
```



# Values Changes in Distributed Environment

- What happens when you perform a commit based on an old value?
  - When the record has changed
- Use pessimistic locks to eject all changes until you have completed your update
  - Bad because will slow the application because of contention





# Conditional Update

- Updating a single record may be incorrect if the original value of the read has changed
- Perform the update based on an expected value of a field
  - Eg a timestamp
  - Compare and set method
- Eg. Withdrawing \$70 from an account
  - Current account balance is \$100
  - Update balance to \$30 provided the balance is still \$100

```
select @old_balance := balance,  
       @last_update := update_time  
from accounts  
where acct_id = 'abc123';
```

```
insert into accounts  
set balance = balance - @amount  
where balance = @old_balance and  
       update_time = @last_update
```



# Implementing Conditional Update


- Using a extra column to hold the timestamp of the last update
  - Column will be automatically updated when the record is updated

```
create table accounts (  
    acct_id varchar(8) not null,  
    balance decimal(10,2),  
    last_update default current_timestamp  
        on update current_timestamp,  
    primary key(acct_id)  
)
```

Update this field with the  
timestamp whenever the record  
is updated

```
select * from accounts  
where acct_id = 'abc123'
```

```
update accounts  
set balance = <new_balance>  
where last_update = <last_update from select>
```





# What are Transactions?

- Group multiple operations into a single atomic unit of work
  - All operations within this group must succeed or not performed at all
  - Keep the database in a consistent state
- At the end of a transaction
  - Commit - save all the work performed inside the transaction
  - Rollback - discard all the work performed inside the transaction
- Transactions are isolated and independent of any other transactions
  - Happening at the same time
- Transactions must exhibit ACID property



# ACID Properties




## Atomicity

Each transaction is all or nothing




## Consistency

Data should be valid according to the all defined rules



## Isolation

Transactions do not affect each other



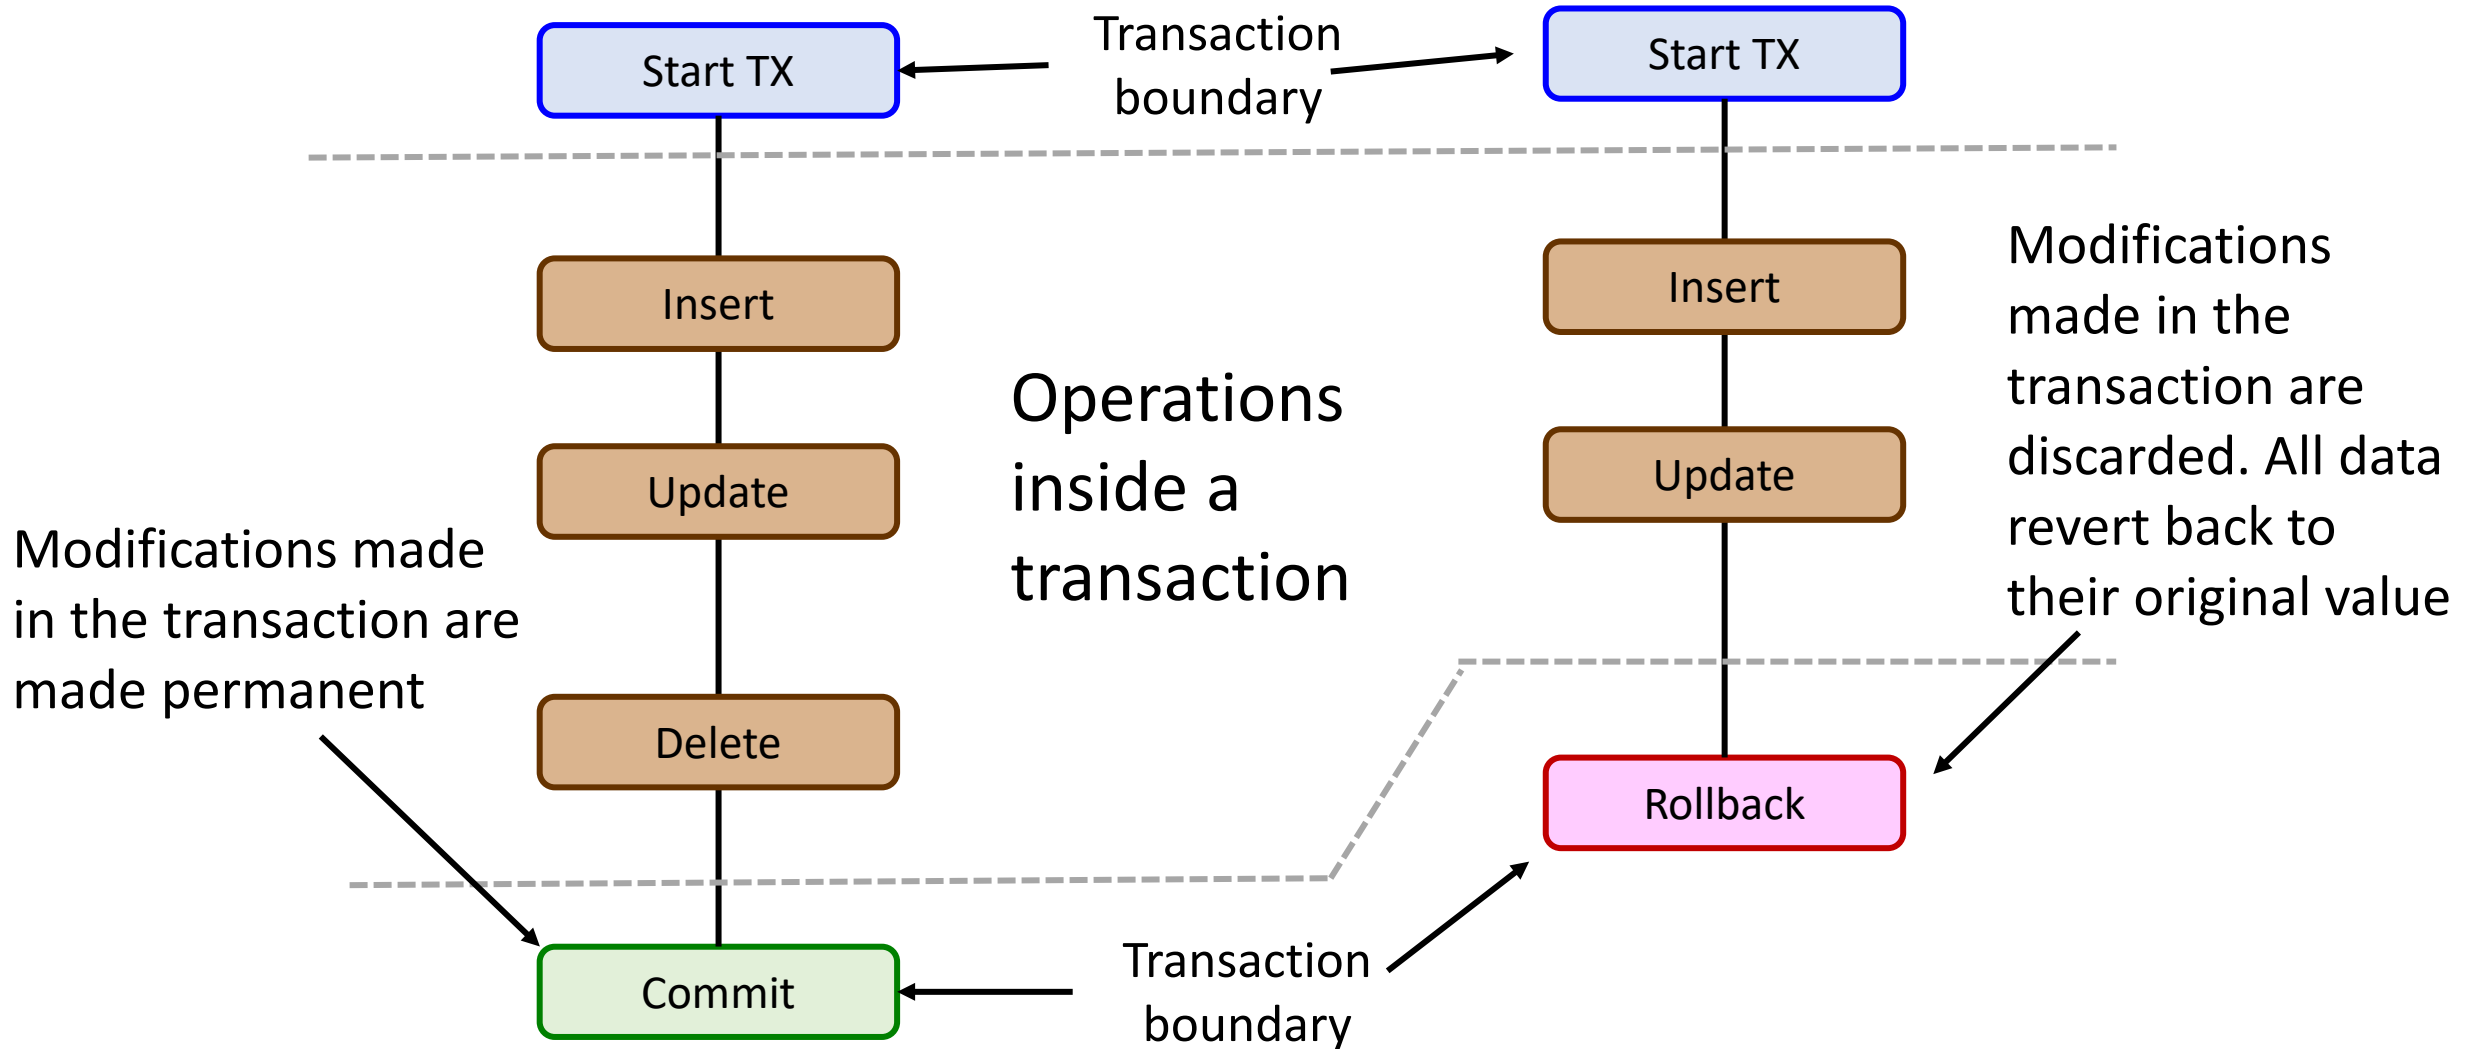
## Durability

Committed data would not be lost






# Transaction





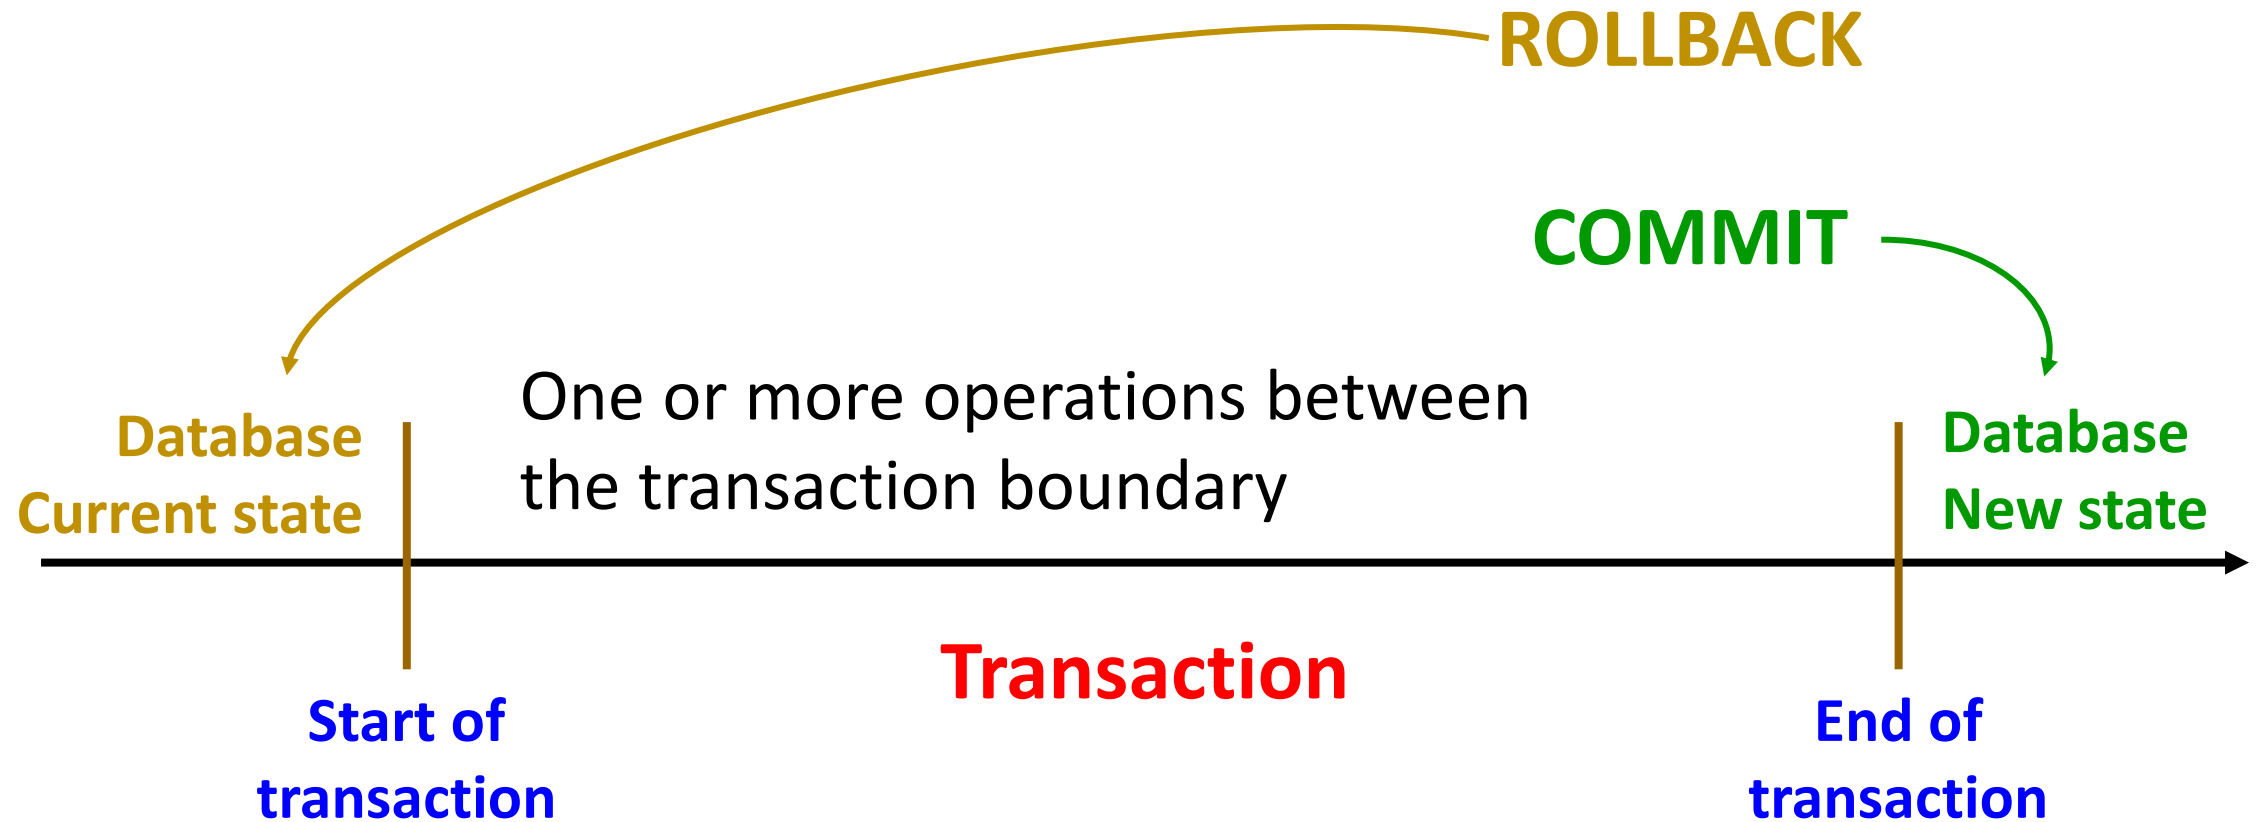
# Transaction

All operations inside the transaction is considered atomic





# Transaction





# When to Use Transaction?

- When you have to modify related data from a few different tables from the same database
- If anyone of the operation fails, then must revert all the modified data back to its original value
  - Eg. Funds transfer from one account to another
- Prevent another operation from interfering with you operation before you commit your data
  - Eg. Prevent others from reading your new address before you finish the update



# Funds Transfer - Atomicity



Fred

Transfer \$100

Barney



Balance: \$500



Balance: \$300

 \$100



Balance: \$400

 \$100



Balance: \$400



# Funds Transfer - Atomicity



Fred

Transfer \$100

Barney



Balance: \$500



Balance: \$300

 \$100



Balance: \$400



Inconsistent



Balance: \$300



# Transaction in SQL

Assume that there are enough funds in `from_acct` for the transfer

```
set @from_acct = 'fred';  
set @to_acct = "barney";  
set @amount = 100;
```


```
start transaction
```

```
update account set balance = balance - @amount  
where acct_id = @from_acct;
```

```
update account set balance = balance + @amount  
where acct_id = @to_acct;
```

```
commit;
```

The 2 update statement must complete successfully or not executed at all







# Performing Transaction

- Methods that are annotated with `@Transaction`
  - Typically methods in `@Service` but may also be present on other bean eg. `@Repository`
- Transaction are automatically committed when the method completes successfully
- Transaction will rollback when an unchecked exception is thrown
  - Exceptions subclass from `RuntimeException`
- Can specify a specific transaction with `rollbackFor` attribute



# Example - Funds Transfer

**begin transaction**

**update account set balance = balance - amount where acct\_id = 'fred'**

**update account set balance = balance + amount where acct\_id = 'barney'**

**commit**



# Example - Funds Transfer

```
@Repository
public class AccountRepository {

    @Autowired private JdbcTemplate template;

    // May throw unchecked exception DataAccessException
    public boolean withdraw(String acctId, double amount) {
        final int rowCount = template.update(
            "update account set balance = balance - ? where acct_id = ?",
            amount, accountId);
        return rowCount > 0;
    }

    public boolean deposit(String acctId, double amount) {
        final int rowCount = template.update(
            "update account set balance = balance + ? where acct_id = ?",
            amount, accountId);
        return rowCount > 0;
    }
}
```



# Example - Funds Transfer

```
@Repository
public class AccountRepository {

    @Autowired private JdbcTemplate template;

    ...

    public Optional<double> getBalance(String acctId) {
        final SqlRowSet rs = template.query(
            'select balance from account where acct_id = ?', acctId);
        return Optional.ofNullable(
            rs.next()? rs.getDouble("balance"): null);
    }
}
```



# Example - Funds Transfer

```
@Service
public class FundsTransfer {
    @Autowired AccountRepository acctRepo;
```



Transaction will rollback when these exceptions are thrown

**@Transactional**

```
public void transfer(String fromAcct, String toAcct, double amount) {

    final Optional<double> optFrom = acctRepo.getBalance(fromAcct);
    final Optional<double> optTo = acctRepo.getBalance(toAcct);
    if (optFrom.isEmpty() || optTo.isEmpty() || (optFrom.get() < amount))
        throw new IllegalArgumentException("Incorrect parameters");

    if !(acctRepo.withdraw(fromAcct, amount) ||
        acctRepo.deposit(toAcct, amount))
        throw new DataAccessException("Cannot perform transfer");

}
```

Transaction boundary

Transaction commits when exit.  
Updates are now permanent





# Isolation

- Multiple transactions can occur independently at the same time without interfering each other
- The intermediate states of an inflight transaction is invisible to other transactions
  - Partially updated records in a fund transfer should not be visible to another transaction or any queries happening at the same time
- Types of isolation
  - Dirty reads – transaction B can read uncommitted data in transaction A
  - Non repeatable reads – a record that is read more than once in a transaction cannot guarantee to return the same value
  - Phantom reads – reading a set of records more than once in a transaction can return more or less records, non repeatable reads at a record level



# Isolation Level

- JDBC drivers use isolation level to provide data integrity during a transaction
  - READ\_UNCOMMITTED
  - READ\_COMMITTED
  - REPEATABLE\_READ
  - SERIALIZABLE
- MySQL JDBC driver default isolation level is set to REPEATABLE\_READ

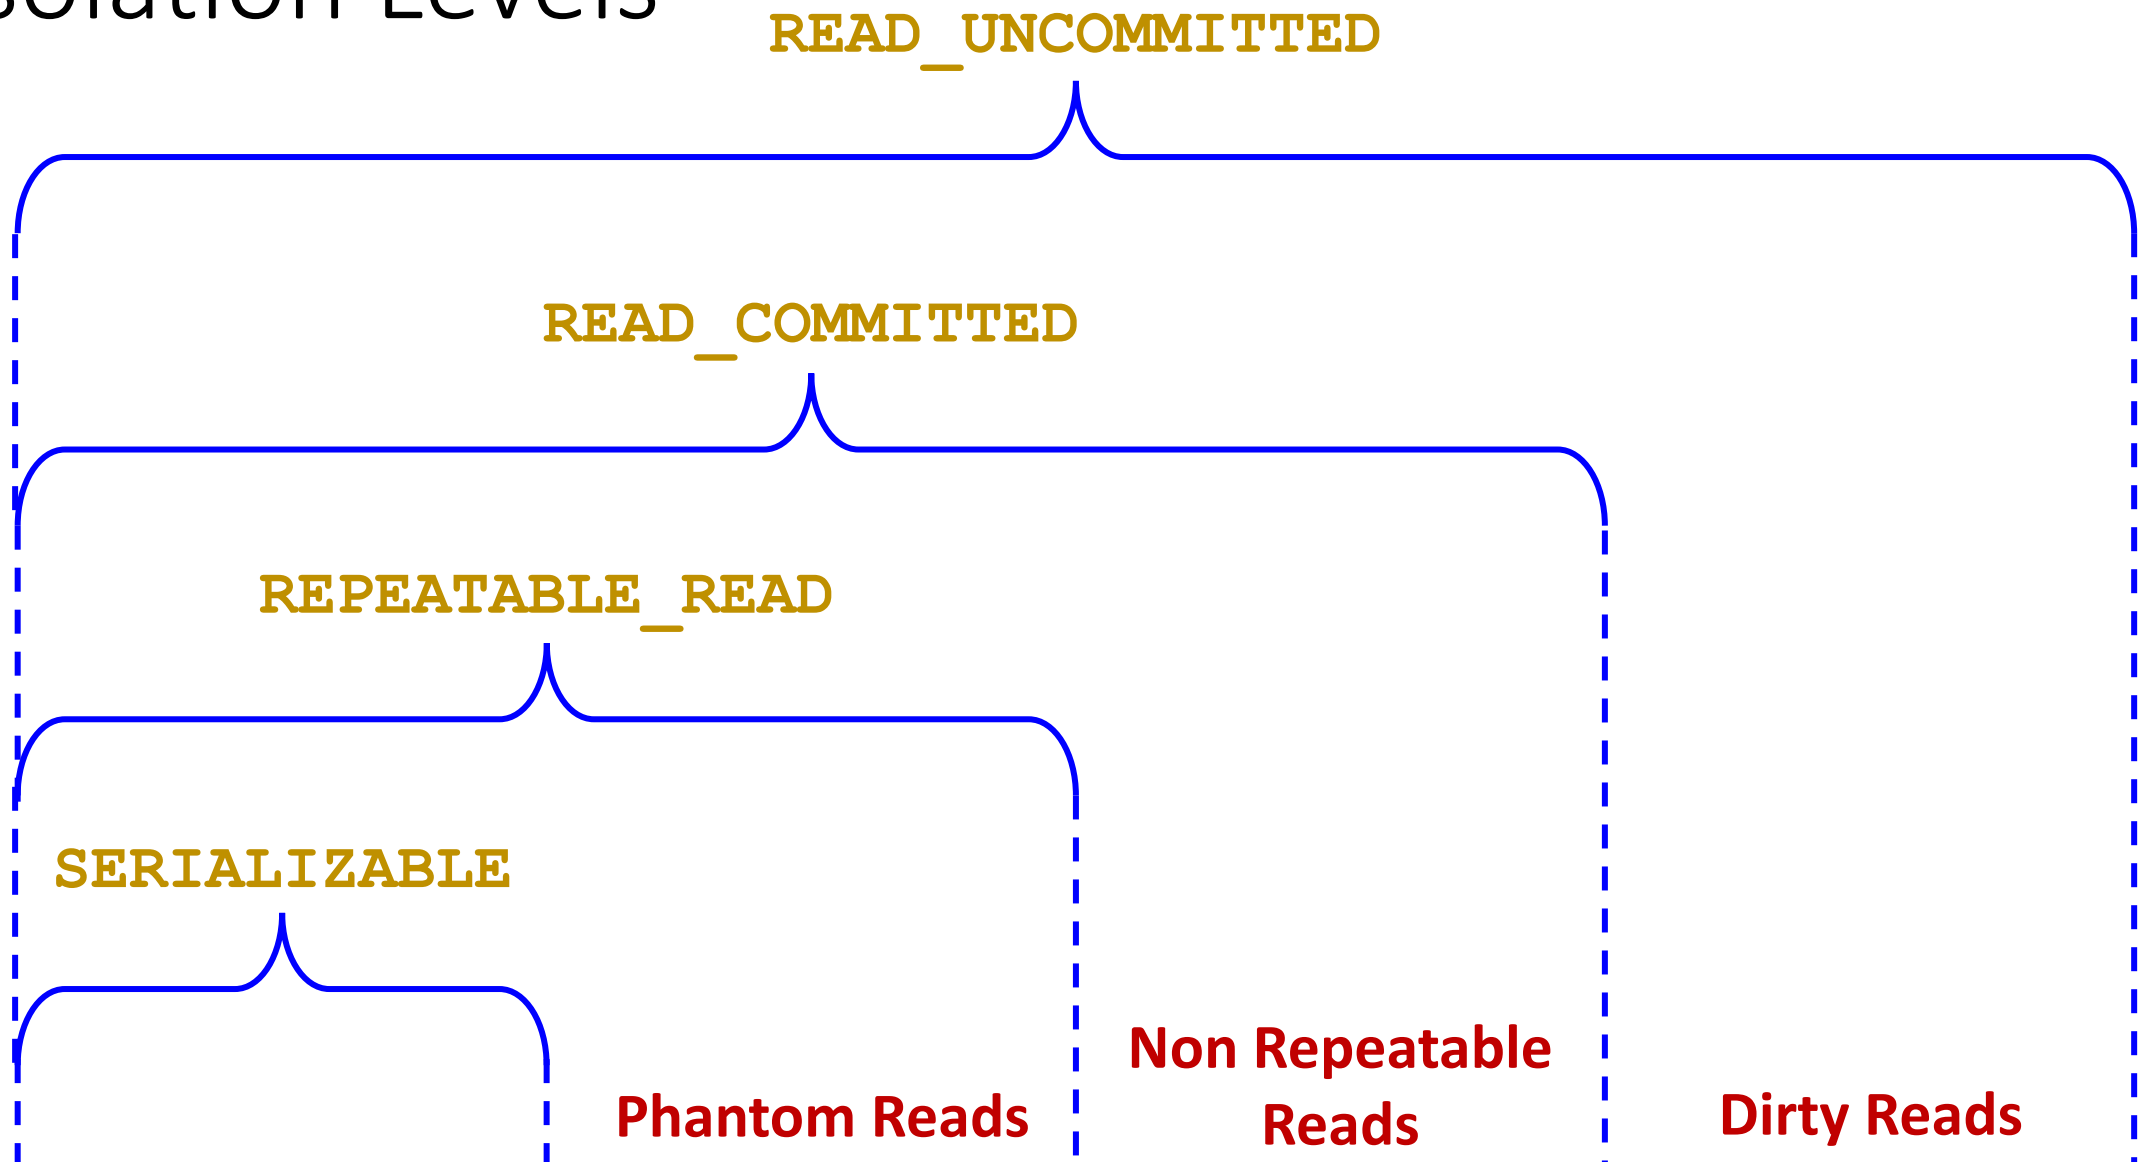
Changing the  
isolation level

```
@Transactional(isolation=Isolation.SERIALIZABLE)
public void transfer(String fromAcct, String toAcct, double amount) {
```





# Isolation Levels



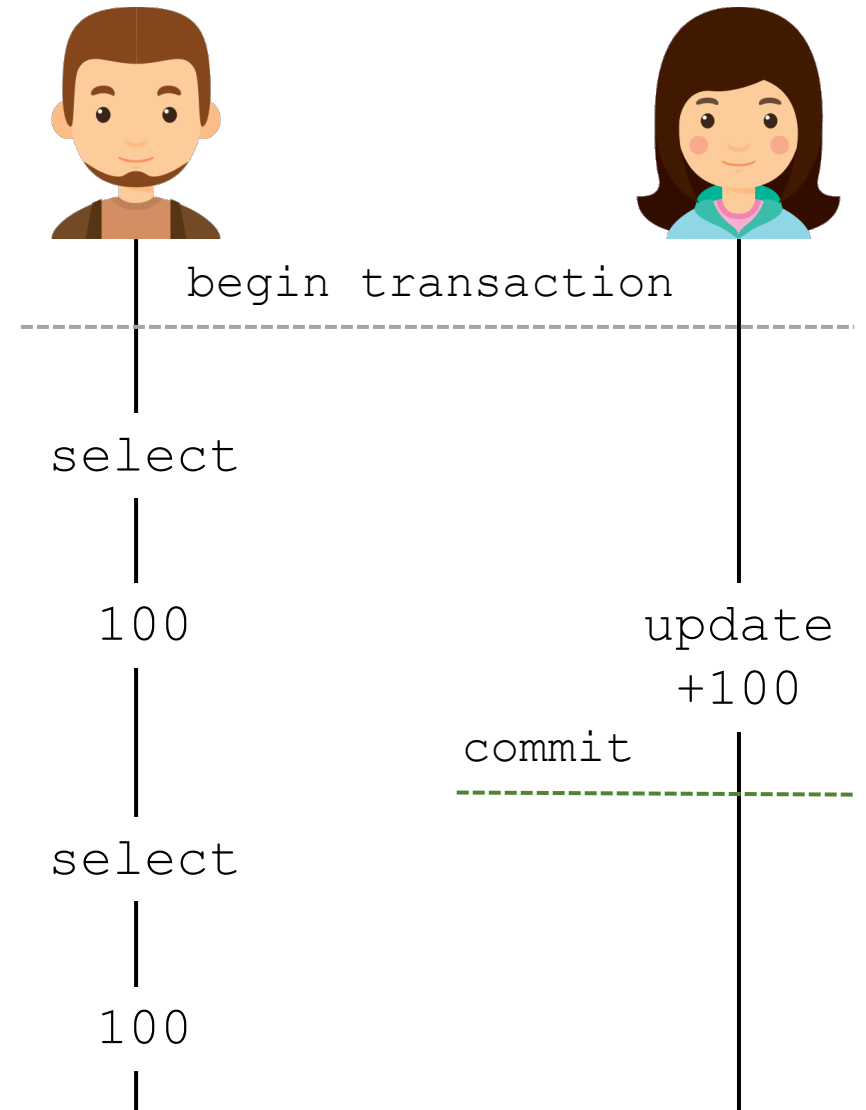


Unused



# Isolation Level - Repeatable Read

- Refers to the visibility of uncommitted results inside a transaction
- Default is Repeatable Reads
- Reads of in flight transactions are isolated from changes
- Defaults to JDBC driver when not set
  - MySQL defaults to repeatable read (TRANSACTION\_READ\_COMMITTED)





# Locks in Transaction

- MySQL will lock the required resources inside a transaction to ensure data consistency
- Row locking for tables with the required index(s)
  - Other sessions can access other rows except for the locked row

```
update account set balance = balance - amount  
where acct_id like 'fred'
```

Primary key has index

- Table locking if the table does not have the appropriate index(s)
  - No session can access any records from the table until transaction completes

```
update account set balance = balance - amount  
where name like 'fred'
```

Non primary key might  
not have index